



Container Engine for Kubernetes (OKE)

Level 200

Jamal Arif
Dec 2018

Safe Harbor Statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

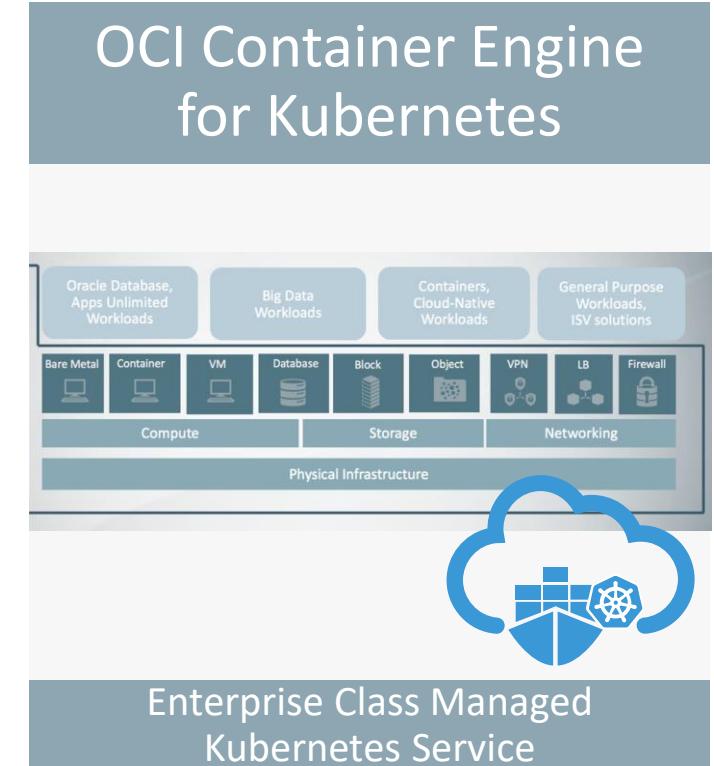
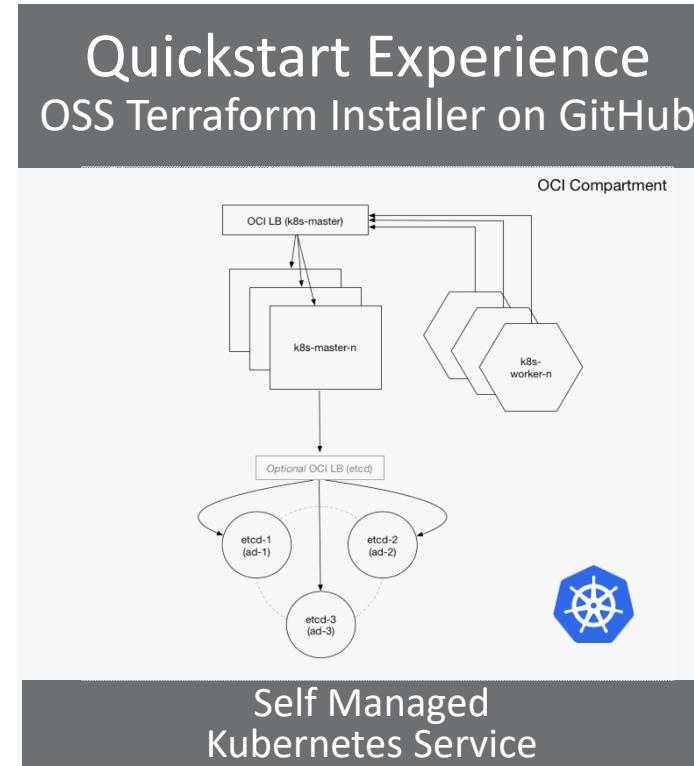
Objectives

After completing this lesson, you should be able to:

- Describe the OCI Container Engine for Kubernetes
- Managing a Kubernetes Cluster on OCI
- **Pre-requisites: Docker and Kubernetes basic understanding**

Oracle Cloud Infrastructure and Kubernetes

Roll Your Own, Pre-Built Installer, Managed Service



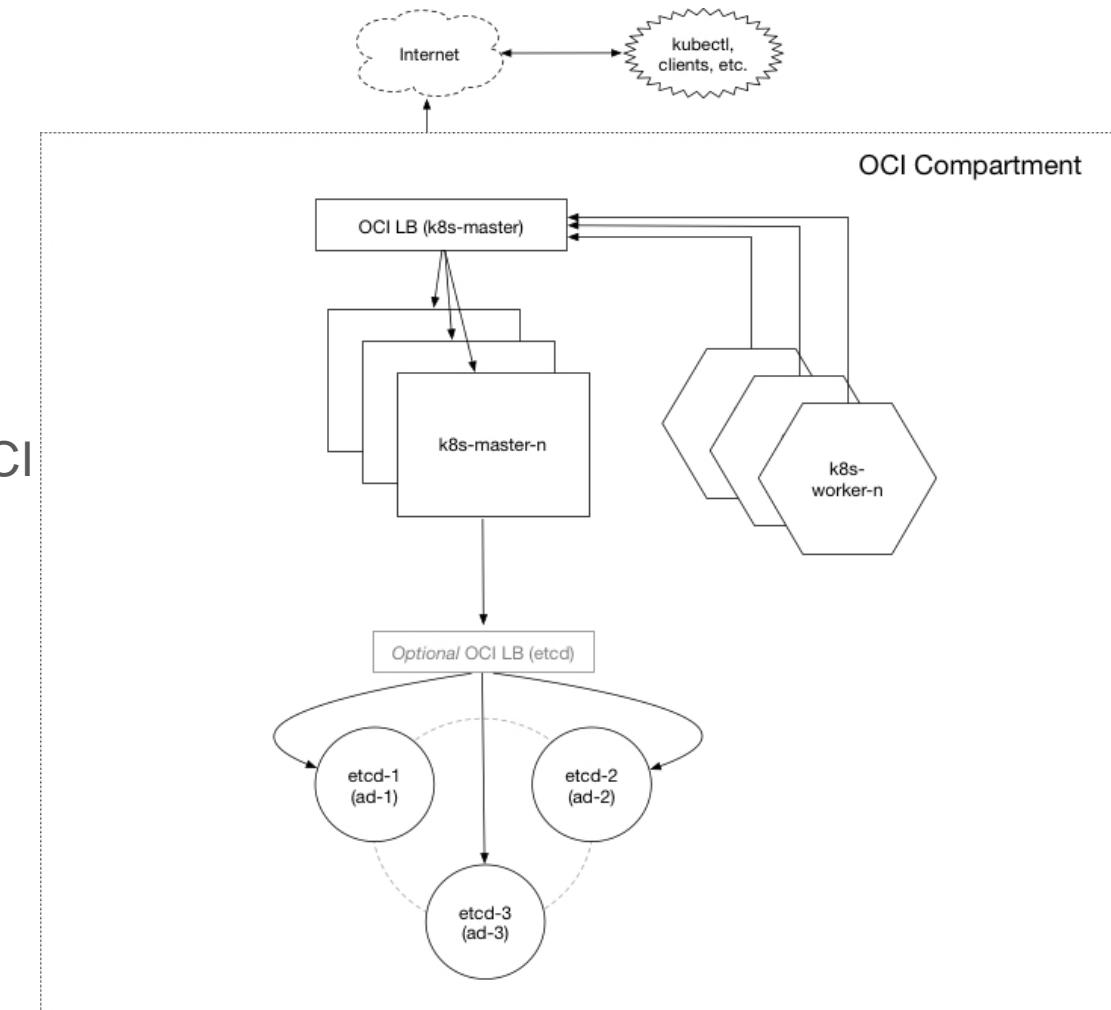
DIY - Terraform Kubernetes Installer for OCI

Open Source OCI Kubernetes installer, based on Terraform

- Oracle developed for Kubernetes on OCI
- Available now on Github -
<https://github.com/oracle/terraform-kubernetes-installer>

Key Highlights

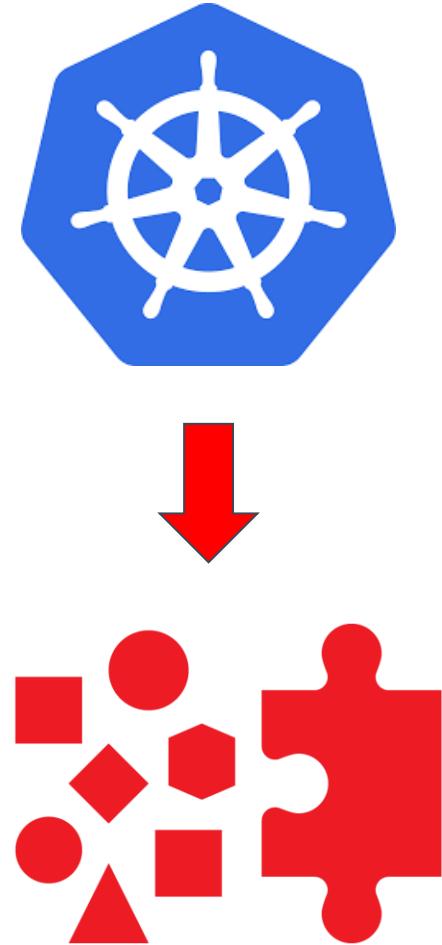
- Highly available Kubernetes cluster configured in your OCI tenancy and compartment
- Creates VCN, subnets, LBs and instances for control plane
- Specify number and shape of nodes for your cluster
- Scale your cluster as needed



[Available on Oracle Github!](https://github.com/oracle/terraform-kubernetes-installer)

Kubernetes Challenges

- Managing, maintaining, upgrading Kubernetes Control Plane
 - API Server, etcd, scheduler etc....
- Managing, maintaining, upgrading Kubernetes Data Plane
 - In place upgrades, deploy parallel cluster etc....
- Figuring out container networking & storage
 - Overlays, persistent storage etc... - it should just work
- Managing Teams
 - How do I manage & control team access to my clusters?
- CI/CD Integration
 - How do I drive automated testing and conditional release into my application lifecycle?





Introducing Container Engine for Kubernetes - OKE

What is It?

- Managed Kubernetes container service to deploy and run your own container based apps
- Tooling to create, scale, manage & control your own standard Kubernetes clusters instantly

What Problems Does it Solve?

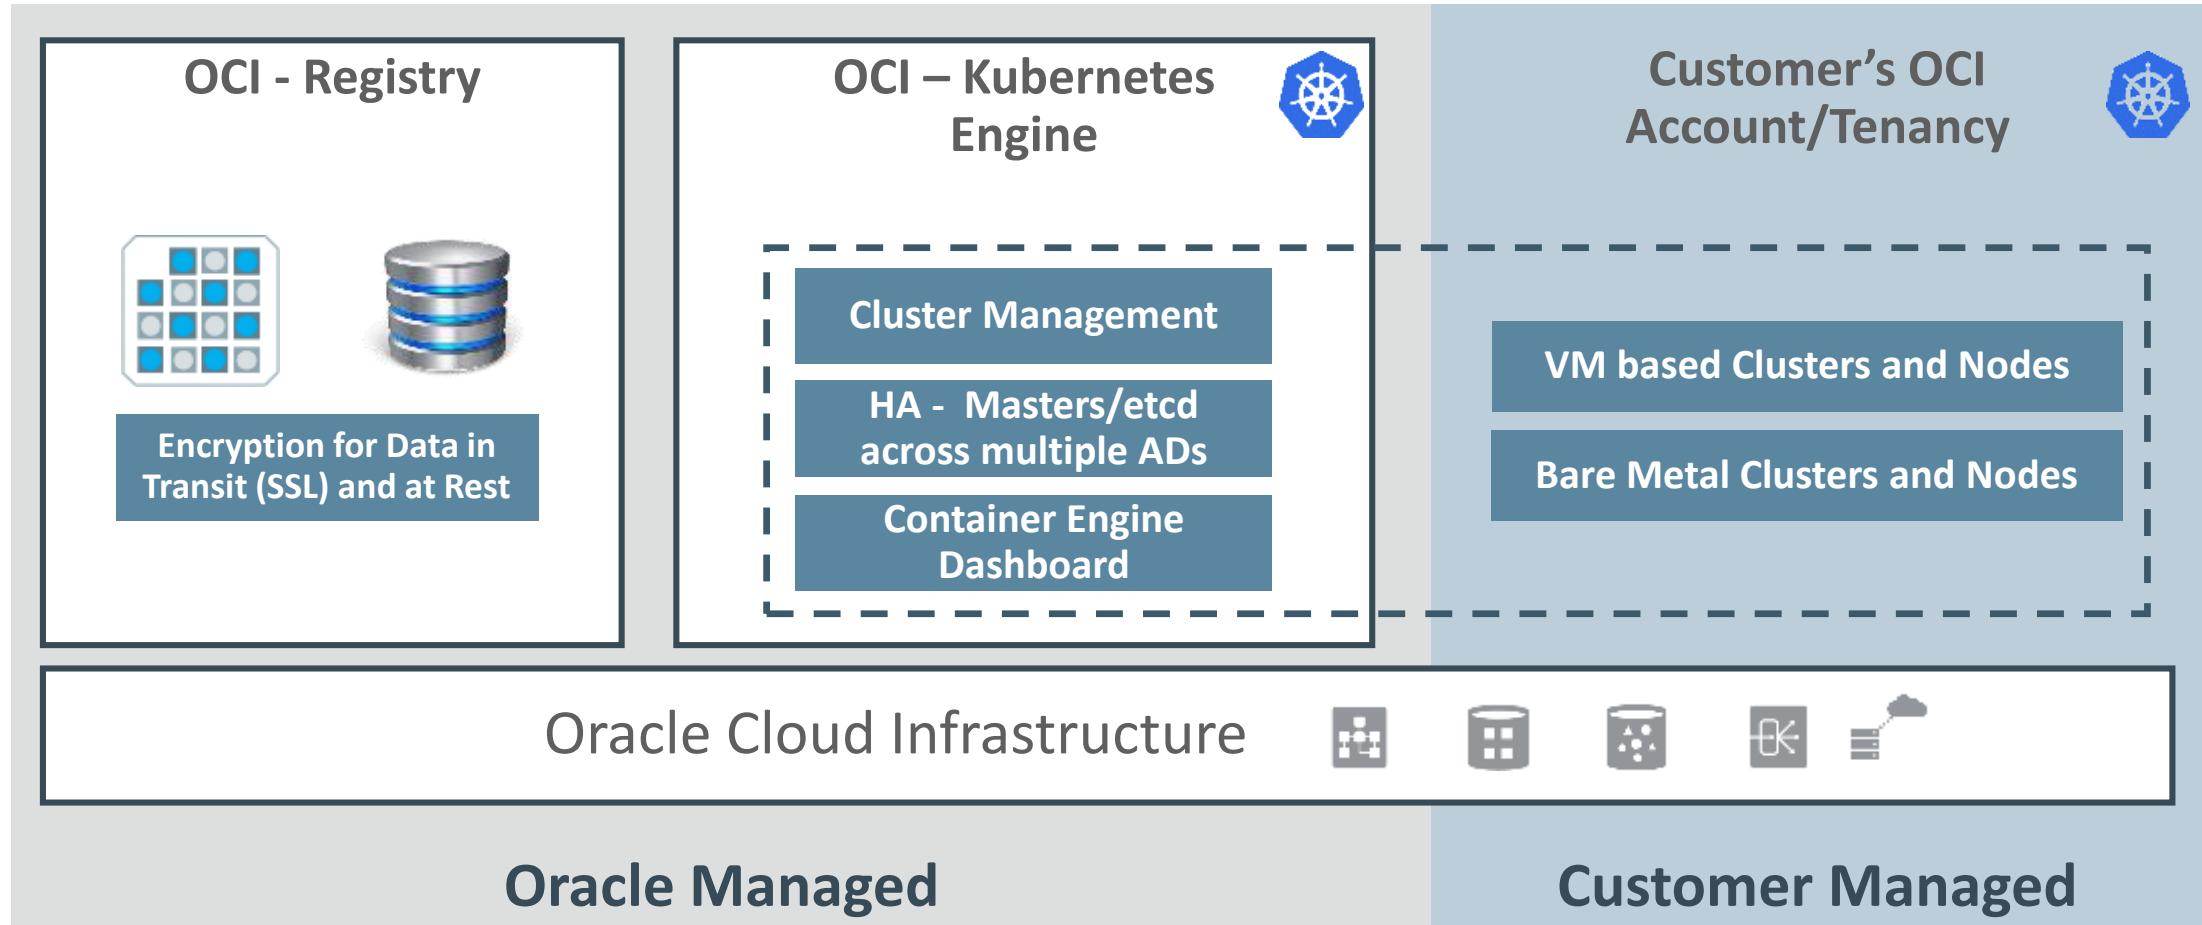
- Too complex, costly and time consuming to build & maintain Kubernetes environments
- Too hard to integrate Kubernetes with a registry and build process for container lifecycle management
- Too difficult to manage and control team access to production clusters

Key Benefits

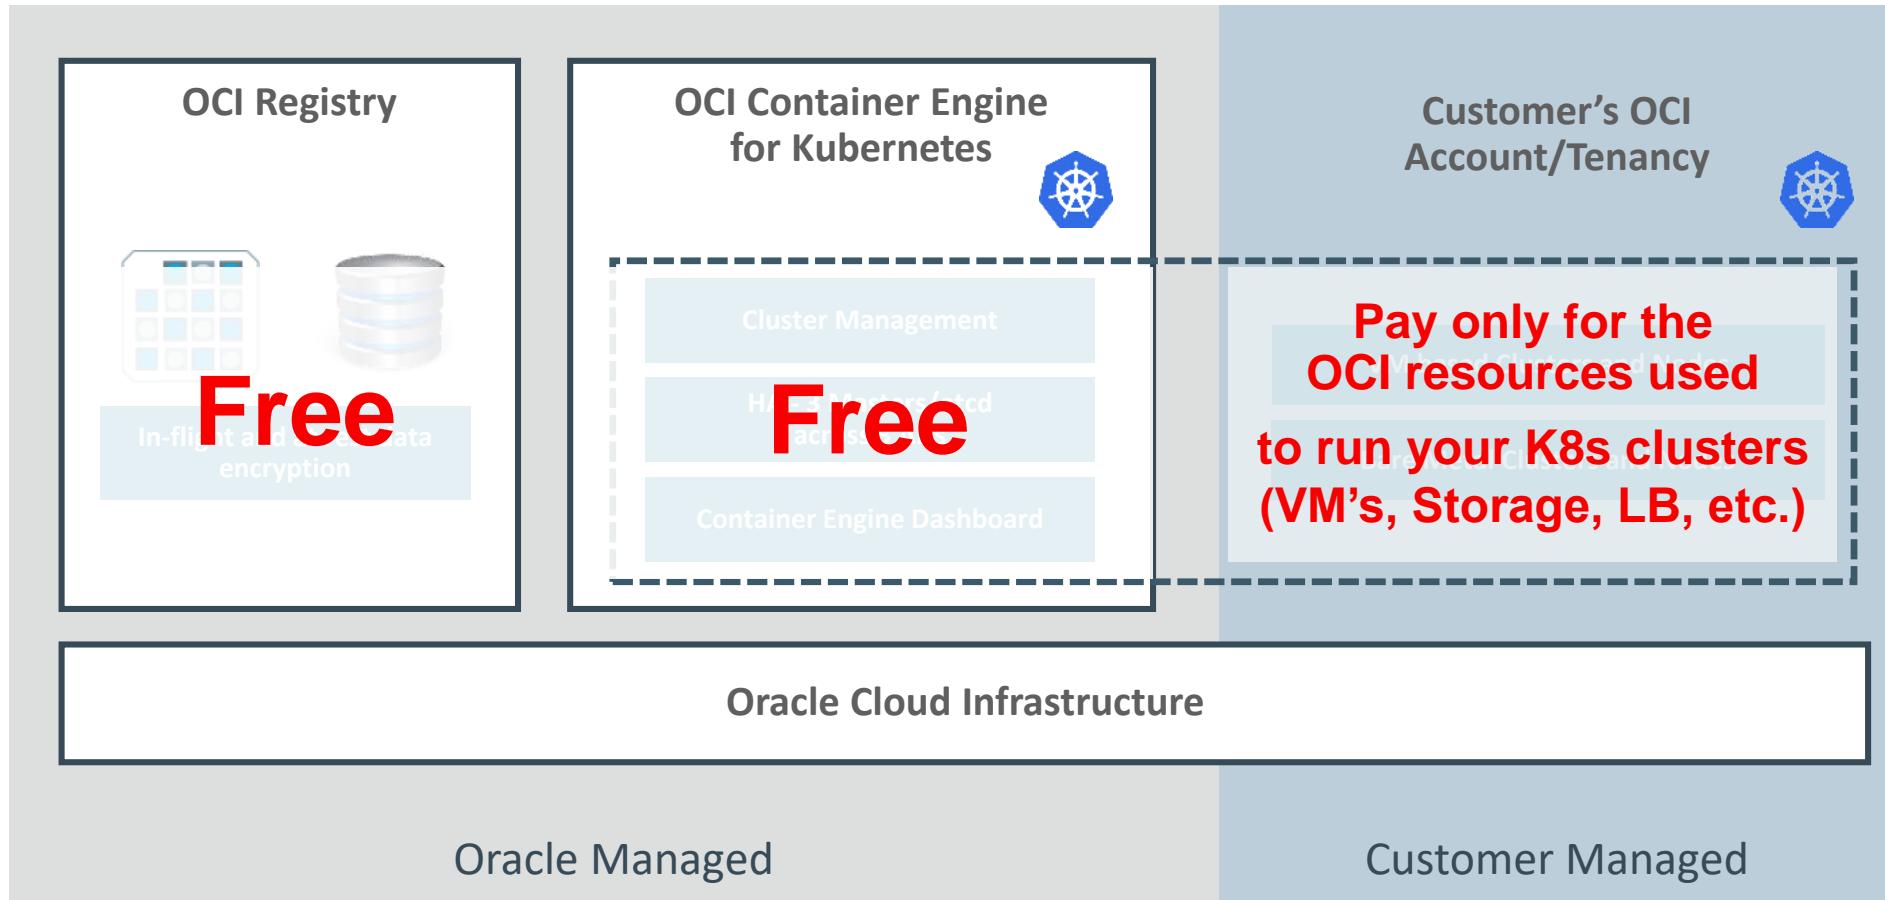
- Enables developers to get started and deploy containers quickly. Gives DevOps teams visibility and control for Kubernetes management.
- Combines production grade container orchestration of open Kubernetes, with control, security, IAM, and high predictable performance of Oracle's next generation cloud infrastructure



Working with OKE and OCIR on OCI



OKE/OCIR Pricing and Packaging

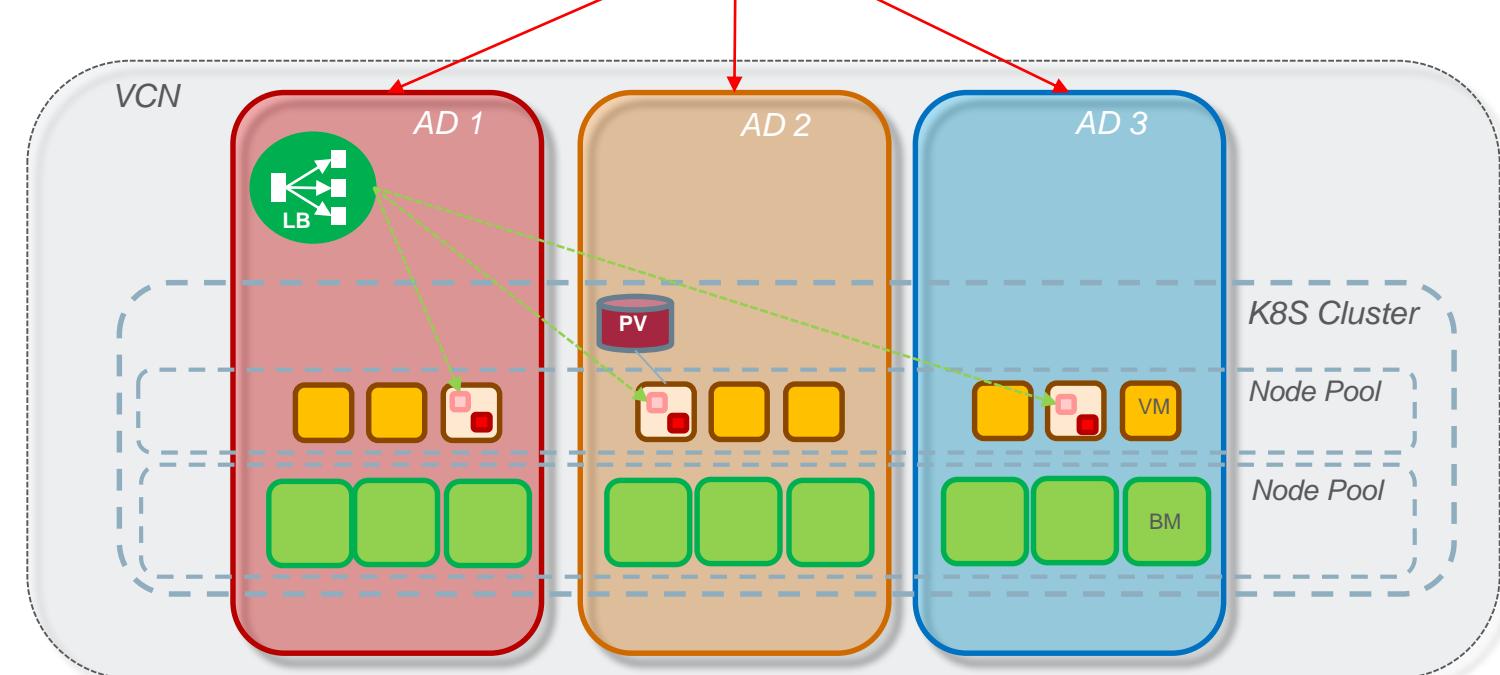
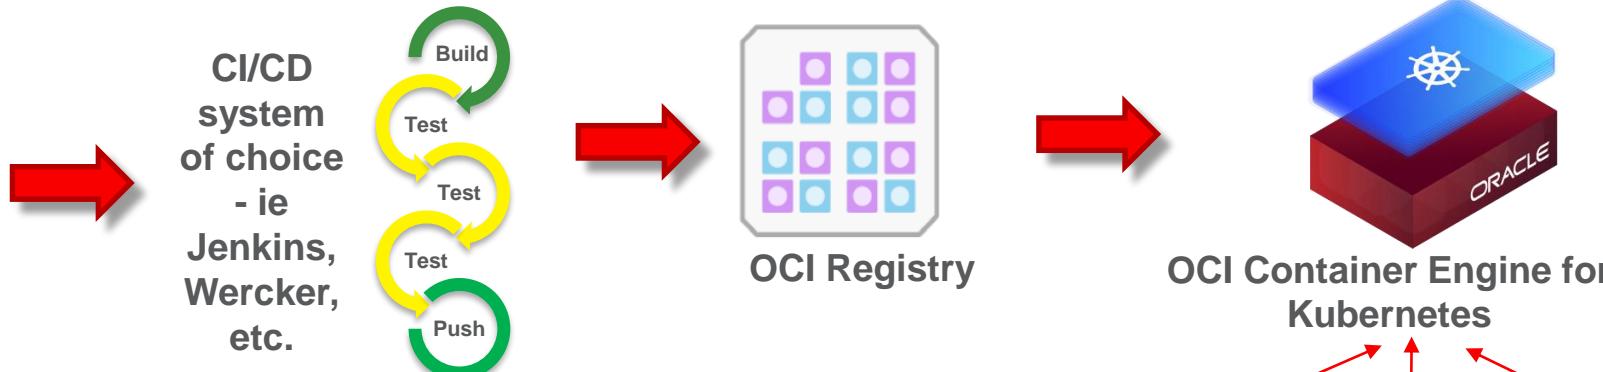


OCI Container Engine for Kubernetes and Registry

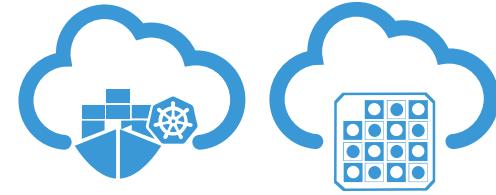
An Open, Fully-Managed Kubernetes Platform & Private Registry



Github
(Source)



Oracle Container Engine & Registry



Container Native

- **Standard Docker & Kubernetes**
 - Deploy standard & open upstream Docker and Kubernetes versions for compatibility across environments
- **Registry Integration**
 - Full Docker v2 compatible private registry to store and manage images
- **Container Engine**
 - Deploy and operate containers and clusters
- **Full integration to cloud networking and storage**
 - Leverage the enterprise class networking, load balancing and persistent storage of Oracle Cloud Infrastructure

Developer Friendly

- **Streamlined Workflow**
 - Use your favorite CI to push containers to the registry, then Kubernetes to deploy to clusters and manage operations
- **Full REST API**
 - Automate the workflow, create and scale clusters through full REST API
- **Built In Cluster Add-Ons**
 - Kubernetes Dashboard, DNS & Helm
- **Open Standards**
 - Docker Based Runtime
 - Worker Node SSH Access
 - Standard Kubernetes

Enterprise Ready

- **Simplified Cluster Operations**
 - Use the fully managed, highly available registry, master nodes and control plane
- **Full Bare Metal Performance and Highly Available IaaS**
 - Combine Kubernetes with bare metal shapes for raw performance
 - Deploy Kubernetes clusters across multiple Availability Domains for resilient applications
- **Team Based Access Controls**
 - Control team access and permissions to clusters
- **Autonomous Clusters**
 - Maintain cluster size and performance in face of node failures and load fluctuations

Pre-requisites for OKE (1) - Service Limits for tenancy

- Monthly universal Credits have limit of 3 clusters per OCI region with 1000 nodes in a cluster
- Pay-as-you-go or Promo accounts to contact support to activate clusters.
- Must also have compute Instance Quota (**Required**) – to launch k8s worker nodes in an AD or across ADs for HA
- Block Volume Quota – Only required if you want to create k8s persistent volumes
- Load Balancer Quota – Only required if you want to distribute traffic between worker nodes

Pre-requisites for OKE (2) – Required IAM policies

- Required Policy in the root compartment of your tenancy
 - allow service OKE to manage all-resources in tenancy
- To launch a K8s cluster, user must be either part of the Admin group or a group to which a policy grants the appropriate Container Engine for Kubernetes permissions.
- Policies can be created for users which are not part of the admin group
- For Example: To enable users in group 'dev-team' to perform any operation on cluster-related resources
→ allow group dev-team to manage cluster-family in tenancy
 - Note: if users will be using the Console to create and update clusters, polices must also grant the dev-team group the Networking permissions VCN_READ and SUBNET_READ.

Pre-requisites for OKE (3) – Basic Virtual Cloud Network Config

- An Existing VCN with following
 - Internet Gateway
 - Route table with default route to IGW
 - K8s worker node subnets – at least three subnets in different ADs for High Availability of Workers
 - LBs Subnets – 2 subnets in different ADs for OCI Public Load Balancer
 - Separate Security Lists for K8s Worker Nodes Subnets and LB Subnets
 - Security Lists for K8s worker Nodes Subnets
 - Stateless ingress and egress rules that allow all traffic between the different worker node subnets
 - stateless ingress and egress rules that allow all traffic between worker node subnets and load balancer subnets
 - ingress rules to allow Container Engine for Kubernetes to access worker nodes on port 22 from 130.35.0.0/16, 138.1.0.0/17, 147.154.0.0/16, 192.29.0.0/16
 - an egress rule that allows all outbound traffic to the internet

Kubernetes Cluster Creation

- **Quick Cluster** – Uses default settings to create a 'quick cluster' with new network resources as required.
- **Custom Cluster** – custom settings to create a 'custom cluster'. This approach gives you the most control over the new cluster.

Kubernetes Cluster Creation – Quick Cluster

- **Name** – Name of the K8s Cluster
- **Version** - The version of Kubernetes to run on the master node of the cluster
- **Quick Create** – Creates following new resources
 - VCN
 - 2 LB Public Subnets
 - 3 Worker Node Subnets
 - Security List with required security list rules
 - Route table with default gateway to IGW
 - Internet Gateway
- **Node Pool** – Shape of the VM and quantity per subnet with your id_rsa.pub key
- **Kubernetes Labels (key, value)** - Nodes added to this node pool will automatically get one or more Kubernetes labels applied, enabling users to target Kubernetes workloads in a specific pool
- **Additional Add-on** – Kubernetes Dashboard, Tiller (helm)

Kubernetes Cluster Creation – Quick Cluster

Cluster Creation [help](#) [close](#)

CLUSTER COMPARTMENT
tutorials

NAME
cluster2

KUBERNETES VERSION
v1.11.1

Kubernetes version installed on your master and worker nodes

QUICK CREATE
Quickly create a cluster with default settings, also creates a dedicated network

CUSTOM CREATE
Create a cluster with custom settings, assumes an existing network

Create Node Pool

NAME: pool1 COMPARTMENT: tutorials KUBERNETES VERSION: v1.11.1 IMAGE: Oracle-Linux-7.5

SHAPE QUANTITY PER SUBNET

VM.Standard2.1 1

The shape of all nodes in the pool
The number of nodes per subnet.

PUBLIC SSH KEY OPTIONAL
Input SSH public key

The SSH public key to access your nodes.

Kubernetes Labels

KEY	VALUE
name	pool1

Nodes added to this node pool will automatically get one or more Kubernetes labels applied, enabling users to target Kubernetes workloads in a specific pool

+ Another Pair

Create Virtual Cloud Network

A new VCN network will be created for you in order to have a functioning cluster

COMPARTMENT: tutorials RESOURCE CREATION: 1 VCN, 2 service lb subnets and 3 worker node subnets

Additional Add Ons

- KUBERNETES DASHBOARD ENABLED
- TILLER (HELM) ENABLED

Kubernetes Cluster Creation – Custom Cluster

- **Name** – Name of the K8s Cluster
- **Version** - The version of Kubernetes to run on the master node of the cluster
- **VCN** - The name of an existing virtual cloud network that has been configured for cluster creation and deployment
- **Kubernetes Service LB Subnets:** The two subnets configured to host load balancers.
- **Kubernetes Service CIDR Block (Optional):** The available group of network addresses that can be exposed as Kubernetes services (ClusterIPs), expressed as a single, contiguous IPv4 CIDR block. For example, 10.96.0.0/16. Must not overlap with VCN CIDR
- **Pods CIDR Block (Optional):** The available group of network addresses that can be allocated to pods running in the cluster, expressed as a single, contiguous IPv4 CIDR block. For example, 10.244.0.0/16. Must not overlap with VCN CIDR
- **Kubernetes Dashboard and Helm** are enabled by default

Kubernetes Worker Nodes – Node pools

- **Name** – Name of the node pool
- **Version** - The version of Kubernetes to run on each worker node in the node pool. By default, the version of Kubernetes specified for the master node is selected. The Kubernetes version on worker nodes must be either the same version as that on the master node, or an earlier version that is still compatible
- **Image:** The image to use on each node in the node pool. An image is a template of a virtual hard drive that determines the operating system and other software for the node.
- **Shape:** The number of CPUs and the amount of memory allocated to each node in the node pool.
- **Subnet:** One or more subnets configured to host worker nodes. The worker node subnets must be different to the load balancer subnets.
- **Quantity per Subnet:** The number of worker nodes to create for the node pool in each subnet.
- **Public SSH Key:** (Optional) The public key portion of the key pair you want to use for SSH access to each node in the node pool.
- **Kubernetes Labels (key, value)** - Nodes added to this node pool will automatically get one or more Kubernetes labels applied, enabling users to target Kubernetes workloads in a specific pool

Kubernetes Cluster Creation – Custom Cluster

Cluster Creation [help](#) [close](#)

CLUSTER COMPARTMENT
tutorials

NAME
 

KUBERNETES VERSION
 
Kubernetes version installed on your master nodes

QUICK CREATE
Quickly create a cluster with default settings, also creates a dedicated network

CUSTOM CREATE
Create a cluster with custom settings, assumes an existing network 

Kubernetes Cluster Creation – Custom Cluster

Network Selection

NETWORK COMPARTMENT

tutorials
jamalarif (root)/tutorials

VCN

oke-vcn-quick-cluster1-20181204205349

An existing VCN to provision your cluster in. A /16 CIDR is sufficient for the majority of cases. An Internet Gateway is required as well as a default route to the gateway.

KUBERNETES SERVICE LB SUBNETS

oke-svclbsubnet-quick-cluster1-20181204205349-fyhg:PHX-AD-2
oke-svclbsubnet-quick-cluster1-20181204205349-fyhg:PHX-AD-1

If automatic Load Balancer integration is desired, two subnets must be provided to host Load Balancers created by Kubernetes in your tenancy. These subnets should be different from those subnets used by node pools that you create for this cluster (cluster nodes).

KUBERNETES SERVICE CIDR BLOCK OPTIONAL

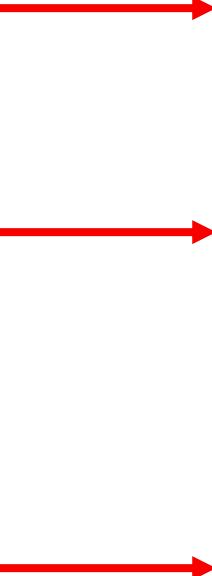
Defaults to 10.96.0.0/16

This is the CIDR range used by exposed Kubernetes services (ClusterIPs). This CIDR should not overlap with the VCN CIDR range.

PODS CIDR BLOCK OPTIONAL

Defaults to 10.244.0.0/16

This is the CIDR range used for IP addresses by your pods. A /16 CIDR is generally sufficient. This CIDR should not overlap with any subnet range in the VCN (it can also be outside the VCN CIDR range).



Additional Add Ons

- KUBERNETES DASHBOARD ENABLED
- TILLER (HELM) ENABLED

Kubernetes Cluster Creation – Custom Cluster

The screenshot shows the 'Node Pool' configuration screen for creating a custom Kubernetes cluster. The left panel contains fields for NAME, VERSION, IMAGE, SHAPE, and SUBNETS. The right panel contains fields for QUANTITY PER SUBNET, PUBLIC SSH KEY, and KUBERNETES LABELS.

Left Panel: Node Pool Configuration

- NAME:** NodePoolCustom (highlighted by a red arrow)
- VERSION:** v1.11.1 (highlighted by a red arrow)
- IMAGE:** Oracle-Linux-7.5 (highlighted by a red arrow)
- SHAPE:** VM.Standard2.1 (highlighted by a red arrow)
- SUBNETS:** oke-subnet-quick-cluster1-20181204205349-fyhg:PHX-AD-3, oke-subnet-quick-cluster1-20181204205349-fyhg:PHX-AD-1, oke-subnet-quick-cluster1-20181204205349-fyhg:PHX-AD-2 (highlighted by a red arrow)

Right Panel: Additional Configuration

- QUANTITY PER SUBNET:** 1 (highlighted by a red arrow)
- PUBLIC SSH KEY (OPTIONAL):** Input SSH public key (The SSH public key to access your nodes.)
- KUBERNETES LABELS:** prod, AppA (highlighted by a red box)

Kubernetes Cluster Details and Node Pools

OKE-Cluster-Demo

[Access Kubeconfig](#) [Delete Cluster](#)

[Cluster Details](#)

Cluster Status: ✓ Active	Kubernetes Version: v1.9.7
Node Pools: 1	Kubernetes Address: ...com:6443 Show Copy
Cluster ID: ...izdemey4d Show Copy	VCN Name: OKE-VCN-Ashburn
Launched: Thu, 17 May 2018 20:07:18 GMT	VCN Id: ...wbqnt7bq Show Copy
Services Cidr: 10.96.0.0/16	Pods Cidr: 10.244.0.0/16

Node Pools

[Add Node Pool](#)

NodePool1				Actions ▾
Kubernetes Ver: v1.9.7	Image Name: Oracle-Linux-7.4	Nodes Per Subnet: 1		
Shape: VM.Standard1.2	Total Worker Nodes: 3	Number of Subnets: 3		
Hide Node Details				
Instance Name ▾	Node State	Subnet	Public IP	
oke-czgizdemey4d-n4dgijwge2g-sbj5bkk37a-0	ACTIVE	Sub-AD1	129.213.84.121	
oke-czgizdemey4d-n4dgijwge2g-sirhjog44iq-0	ACTIVE	Sub-AD2	129.213.57.201	
oke-czgizdemey4d-n4dgijwge2g-sw5634ywka-0	ACTIVE	Sub-AD3	129.213.35.90	
<small>Showing 3 item(s)</small>				

Scaling Node Pools

Node Pools

Add Node Pool

NodePool1

Kubernetes Ver: v1.9.7

Shape: VM.Standard1.2

Image Name: Oracle-Linux-7.4

Total Worker Nodes: 3

Nodes Per Subnet: 1

Number of Subnets: 3

[Hide Node Details](#)

Instance Name ▾

[oke-czgizdemy4d-n4dgyjwge2g-sbpj5bkk37a-0](#)

[oke-czgizdemy4d-n4dgyjwge2g-sirhjog44iq-0](#)

[oke-czgizdemy4d-n4dgyjwge2g-sw56r34ywkg-0](#)

Node State

ACTIVE

Subnet

Sub-AD1

Public IP

129.213.84.121

Showing 3 Item(s)

Actions ▾

Edit

Scale

Delete Node Pool



Scale Node Pool

[help](#) [close](#)

You may scale certain characteristics of a node pool.

SUBNETS

Sub-AD3 Sub-AD2 Sub-AD1

The subnets used for node instances in the node pool. These subnets should be different from the Cluster Kubernetes Service LB subnets. A subnet per Availability Domain (AD) is typical.

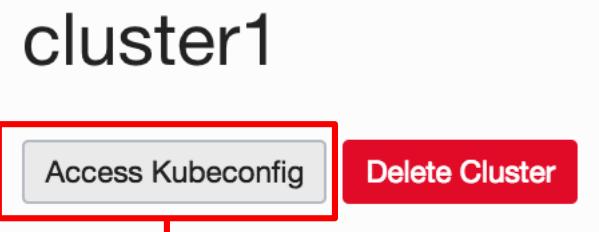
QUANTITY PER SUBNET

1

The number of nodes per subnet.

Scale

Accessing Kubernetes Cluster using kubectl



How to Access Kubeconfig [help](#) [close](#)

You must have [downloaded and installed](#) the OCI CLI and [configured](#) it for use.

To access the kubeconfig for your cluster, run the following commands:

```
1. mkdir -p $HOME/.kube
2. oci ce cluster create-kubeconfig --cluster-id
   ocid1.cluster.oc1.phx.aaaaaaaaaae4wcy3dme3dsmdbmq2wky3emq3wkylgrstonzzhcstomrvgzsw --file
   $HOME/.kube/config --region us-phoenix-1
```

To set your KUBECONFIG environment variable to the file for this cluster, use:

```
export KUBECONFIG=$HOME/.kube/config
```

You may have to add this to your shell initiation script if you wish to persist this change. For more information on managing kubeconfig files, please refer to the official [Kubernetes documentation](#).

More information on the available commands for OCI's Container Engine for Kubernetes CLI can be found [here](#).

[Close](#)

Accessing Kubernetes Cluster using kubectl

The screenshot shows the Oracle Cloud Kubernetes Service console interface. On the left, there's a sidebar with navigation links: Resources, Node Pools, Work Requests, and Getting Started. The 'Getting Started' link is highlighted with a red box and a red arrow points from it to the main content area. The main content area has a title 'Getting Started' and a sub-section 'Kubernetes Dashboard'. It explains that the dashboard provides an overview of applications running in the cluster and information on Kubernetes resources. It lists steps 1. `kubectl proxy` and 2. Dashboard will be available at <http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/>. Below this, there's a 'Quick Start: Deploy Sample App' section with three numbered steps: 1. Access Kubeconfig File, 2. Check Version, and 3. Deploy Application. Each step has associated text and a command-line example.

Getting Started

Kubernetes Dashboard

You can use the Kubernetes Dashboard to get an overview of applications running in your cluster. It also provides information on the state of Kubernetes resources in your clusters, and on any errors that may have occurred.

1. `kubectl proxy`
2. Dashboard will be available at:
<http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/>

Quick Start: Deploy Sample App

- 1 Access Kubeconfig File**

To get started, learn how to download the `kubeconfig` file for this cluster by clicking below. This file will contain a series of authentication mechanisms and cluster connection information.

[Access Kubeconfig](#)
- 2 Check Version**

Verify that kubernetes is available by entering the following command in your terminal

 1. `kubectl version`
- 3 Deploy Application**

Deploy a sample hello world application by running the following command in your terminal.

 1. `kubectl create -f https://k8s.io/docs/tasks/run-application/deployment.yaml`

Kubernetes Dashboard

 **kubernetes** Search + CREATE | ⚙

☰ Overview

Cluster

- Namespaces
- Nodes
- Persistent Volumes
- Roles
- Storage Classes

Namespace

- default ▼

Overview

Workloads

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

Discovery and Load Balancer

- Ingresses
- Services

Config and Storage

- Config Maps
- Persistent Volume Claim
- Secrets

Settings

- About

Workloads

Workloads Statuses

- Deployments: 100.00%
- Pods: 100.00%
- Replica Sets: 100.00%

Deployments

Name	Labels	Pods	Age	Images
<input checked="" type="checkbox"/> nginx	run:nginx	6 / 6	9 days	nginx

Pods

Name	Node	Status	Restarts	Age
<input checked="" type="checkbox"/> nginx-7587c6fdb6-w2kdc	129.213.59.91	Running	0	3 days
<input checked="" type="checkbox"/> nginx-7587c6fdb6-wcpzb	129.213.59.91	Running	0	3 days
<input checked="" type="checkbox"/> nginx-7587c6fdb6-7js72	129.213.59.91	Running	0	9 days
<input checked="" type="checkbox"/> nginx-7587c6fdb6-kzn7b	129.213.59.91	Running	0	9 days
<input checked="" type="checkbox"/> nginx-7587c6fdb6-m4g5f	129.213.59.91	Running	0	9 days
<input checked="" type="checkbox"/> nginx-7587c6fdb6-q24mj	129.213.59.91	Running	0	9 days

Replica Sets

Name	Labels	Pods	Age	Images
<input checked="" type="checkbox"/> nginx-7587c6fdb6	pod-template-hash: 3143729862 run: nginx	6 / 6	9 days	nginx

Discovery and Load Balancing

Services

Name	Labels	Cluster IP	Internal endpoints	External endpoints	Age
<input checked="" type="checkbox"/> kubernetes	component: apiserver provider: kubernetes	10.96.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	9 days

Upgrading Kubernetes Master

- OKE Service supports **in-place** upgrade of Kubernetes Master nodes (via Console or API)
 - After upgrading a master node to a newer version of Kubernetes, you cannot downgrade the master node to an earlier Kubernetes version.
 - The versions of Kubernetes running on the master node and the worker nodes must be compatible (that is, the Kubernetes version on the master node must be no more than two minor versions ahead of the Kubernetes version on the worker nodes)

Containers » Clusters » OKE-Shared-Cluster

OKE-Shared-Cluster

ACTIVE

CL

Access Kubeconfig Delete Cluster Upgrade Available

Cluster Details

Cluster Status: ✓ Active	Kubernetes Version: v1.9.4
Node Pools: 1	Kubernetes Address: ...com:6443 Show Copy
Cluster Id: ...cytdgm4d Show Copy	VCN Name: OKE-VCN
Launched: Sat, 28 Apr 2018 05:31:11 GMT	VCN Id: ...4jewygia Show Copy
Services CIDR: 10.96.0.0/16	Pods CIDR: 10.244.0.0/16

Upgrading Kubernetes Worker Nodes

- Kubernetes Worker nodes are upgraded by performing an ‘out-of-place’ upgrade.
 - To upgrade the version of Kubernetes running on worker nodes in a node pool, you replace the original node pool with a new node pool that has new worker nodes running the appropriate Kubernetes version.
 - ‘drain’ existing worker nodes in the original node pool to prevent new pods starting and to delete existing pods. Once no pods exist, the old node pool can be deleted

Summary

- Describe the OCI Container Engine for Kubernetes
- Launch a Kubernetes Cluster on OCI

ORACLE®
Cloud Infrastructure

cloud.oracle.com/iaas

cloud.oracle.com/tryit